

1. Deploy a Kubernetes cluster for 3 nodes
2. Create a NGINX deployment of 3 replicas

1. Deploy a Kubernetes cluster for 3 nodes

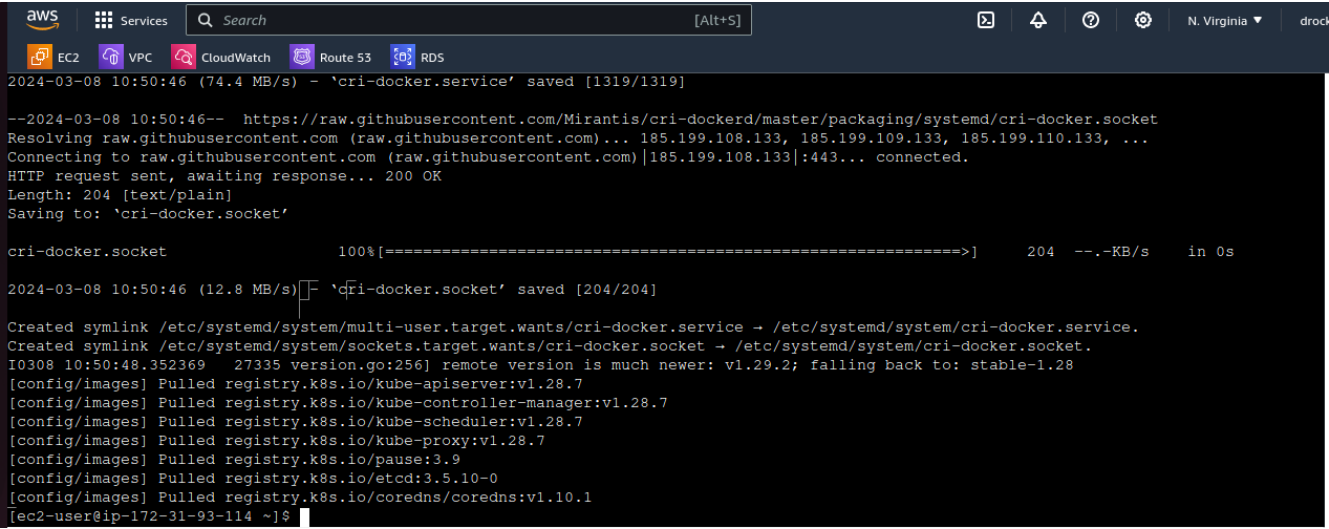
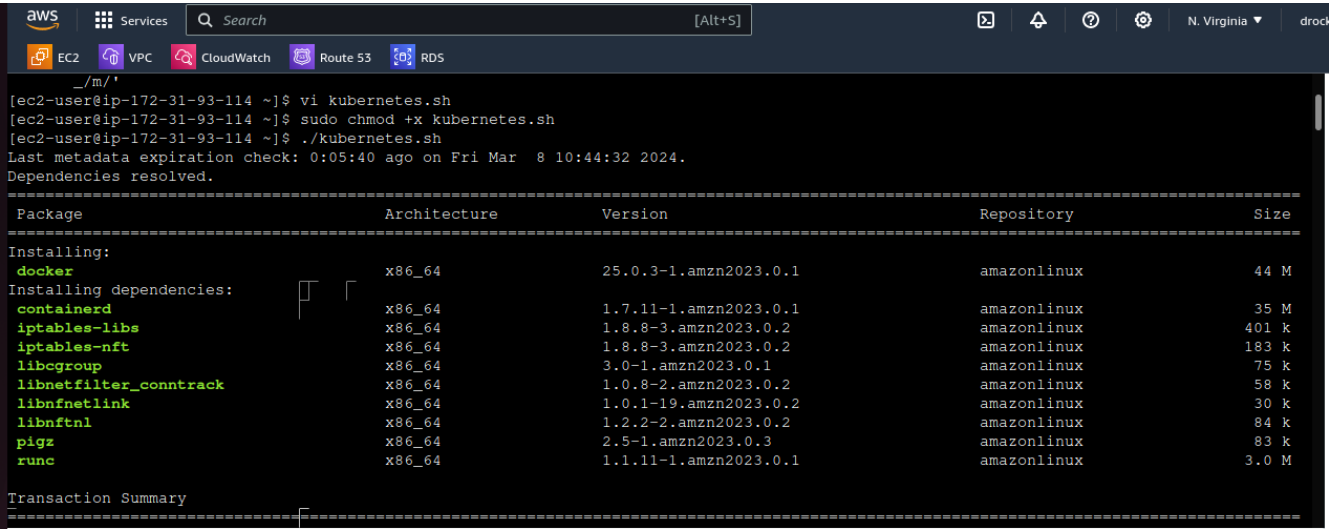
INSTALL 3 LINUX SERVERS; USING minimum of t2-Medium / 20gb hard disk

On all the nodes install all the prerequisites for kubernetes cluster creation: kubernetes.sh

vi kubernetes.sh

```
sudo yum install -y docker
sudo systemctl enable docker
sudo systemctl start docker
sudo setenforce 0
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.28/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.28/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
sudo systemctl enable --now kubelet
VER=$(curl -s https://api.github.com/repos/Mirantis/cri-dockerd/releases/latest|grep tag_name | cut -d
"" -f 4|sed 's/v//g')
wget https://github.com/Mirantis/cri-dockerd/releases/download/v${VER}/cri-dockerd-${
VER}.amd64.tgz
tar xvf cri-dockerd-${VER}.amd64.tgz
sudo mv cri-dockerd/cri-dockerd /usr/local/bin/
wget https://raw.githubusercontent.com/Mirantis/cri-dockerd/master/packaging/systemd/cri-
docker.service
wget https://raw.githubusercontent.com/Mirantis/cri-dockerd/master/packaging/systemd/cri-
docker.socket
sudo mv cri-docker.socket cri-docker.service /etc/systemd/system/
sudo sed -i -e 's,/usr/bin/cri-dockerd,/usr/local/bin/cri-dockerd,' /etc/systemd/system/cri-docker.service
sudo systemctl daemon-reload
sudo systemctl enable cri-docker.service
sudo systemctl enable --now cri-docker.socket
sudo kubeadm config images pull --cri-socket unix:///var/run/cri-dockerd.sock
```

```
chmod +x kubernetes.sh
./kubernetes.sh
```

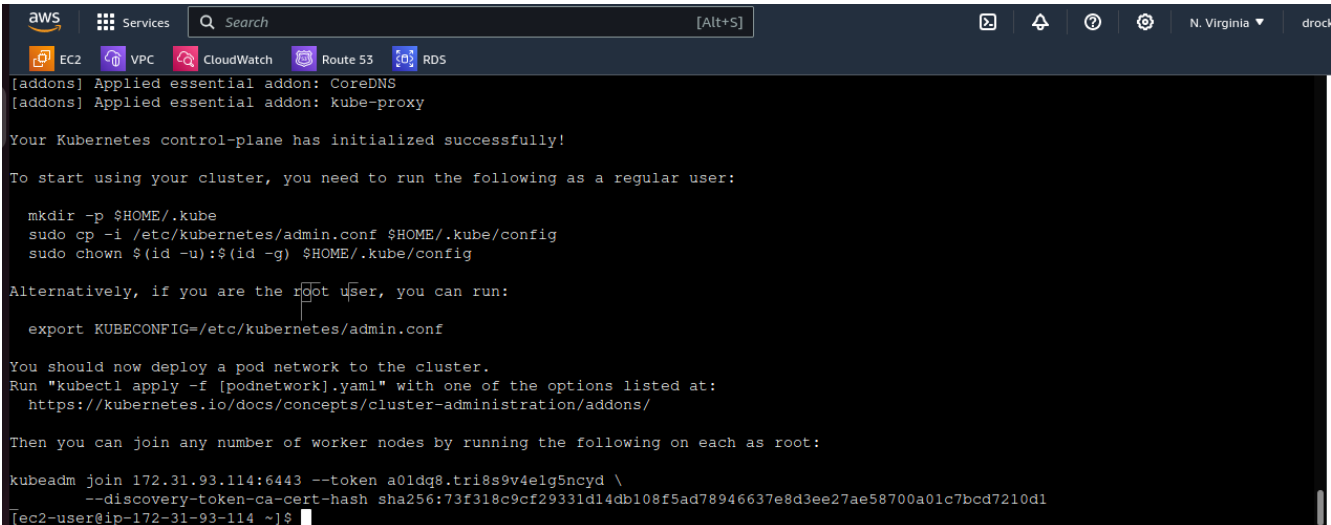


Run the below command on the master server.

```
sudo kubeadm init --pod-network-cidr=192.168.0.0/16 --apiserver-advertise-address=172.31.39.26(Change it to the private IP of ur Master Server) --cri-socket unix:///var/run/cri-dockerd.sock
```

I.e

```
sudo kubeadm init --pod-network-cidr=192.168.0.0/16 --apiserver-advertise-address=172.31.93.114 --cri-socket unix:///var/run/cri-dockerd.sock
```



```
aws Services Search [Alt+S]
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.93.114:6443 --token a01dq8.tri8s9v4e1g5ncyd \
--discovery-token-ca-cert-hash sha256:73f318c9cf29331d14db108f5ad78946637e8d3ee27ae58700a01c7bcd7210d1
[ec2-user@ip-172-31-93-114 ~]$
```

join the worker node in the cluster with this flag

```
--cri-socket unix:///var/run/cri-dockerd.sock
```

Example:

```
sudo kubeadm join 172.31.93.114:6443 --token a01dq8.tri8s9v4e1g5ncyd \
--discovery-token-ca-cert-hash sha256:73f318c9cf29331d14db108f5ad78946637e8d3ee27ae58700a01c7bcd7210d1 --cri-socket
unix:///var/run/cri-dockerd.sock
```

```

[ec2-user@ip-172-31-91-218 ~]$ sudo kubeadm join 172.31.93.114:6443 --token a0ldq8.tri8s9v4e1g5ncyd \
--discovery-token-ca-cert-hash sha256:73f318c9cf29331d14db108f5ad78946637e8d3ee27ae58700a01c7bcd7210d1 --cri-socket unix:///var/r
un/cni-dockerd.sock
[preflight] Running pre-flight checks
[WARNING FileExisting-tc]: tc not found in system path
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

This node has joined the cluster:
* Certificate signing request was sent to apiservert and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
[ec2-user@ip-172-31-91-218 ~]$

```

I-0e85835b18a623c9c (kub-slave1)

PublicIPs: 44.201.153.44 PrivateIPs: 172.31.91.218

TO GET RID OF SUDO

```
cd /etc/kubernetes
```

```
ls -al
```

```
cd ~
```

```
mkdir -p $HOME/.kube
```

```
cd .kube
```

sudo cp -i /etc/kubernetes/admin.conf \$HOME/.kube/configCOPY THE admin.conf to current user path

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

kubectl get nodes ,,,,,,,,,,NOT READY ,

```

[ec2-user@ip-172-31-93-114 ~]$ cd /etc/kubernetes
[ec2-user@ip-172-31-93-114 kubernetes]$ ls -la
total 60
drwxr-xr-x. 4 root root 125 Mar 8 11:01 .
drwxr-xr-x. 83 root root 16384 Mar 8 10:50 ..
-rw-r-----. 1 root root 5649 Mar 8 11:01 admin.conf
-rw-r-----. 1 root root 5677 Mar 8 11:01 controller-manager.conf
-rw-r-----. 1 root root 2061 Mar 8 11:01 kubelet.conf
drwxr-xr-x. 2 root root 113 Mar 8 11:01 manifests
drwxr-xr-x. 3 root root 16384 Mar 8 11:01 pki
-rw-r-----. 1 root root 5625 Mar 8 11:01 scheduler.conf
[ec2-user@ip-172-31-93-114 kubernetes]$ mkdir -p $HOME/.kube
[ec2-user@ip-172-31-93-114 kubernetes]$ cd
[ec2-user@ip-172-31-93-114 ~]$ mkdir -p $HOME/.kube
[ec2-user@ip-172-31-93-114 ~]$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
[ec2-user@ip-172-31-93-114 ~]$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
[ec2-user@ip-172-31-93-114 ~]$ kubectl get nodes
NAME                                STATUS    ROLES    AGE    VERSION
ip-172-31-81-175.ec2.internal       NotReady <none>   5m59s  v1.28.7
ip-172-31-91-218.ec2.internal       NotReady <none>   6m32s  v1.28.7
ip-172-31-93-114.ec2.internal       NotReady control-plane 12m    v1.28.7
[ec2-user@ip-172-31-93-114 ~]$

```

CALICO NETWORK CONFIGURE

pass the 3 commands:

```
kubectl create -f https://raw.githubusercontent.com/projectcalico/calico/v3.27.0/manifests/tigera-operator.yaml
```

```
curl https://raw.githubusercontent.com/projectcalico/calico/v3.27.0/manifests/custom-resources.yaml -O
```

```
kubectl create -f custom-resources.yaml
```

```
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippreservations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/apiservers.operator.tigera.io created
customresourcedefinition.apiextensions.k8s.io/imagesets.operator.tigera.io created
customresourcedefinition.apiextensions.k8s.io/installations.operator.tigera.io created
customresourcedefinition.apiextensions.k8s.io/tigerastatuses.operator.tigera.io created
serviceaccount/tigera-operator created
clusterrole.rbac.authorization.k8s.io/tigera-operator created
clusterrolebinding.rbac.authorization.k8s.io/tigera-operator created
deployment.apps/tigera-operator created
[ec2-user@ip-172-31-93-114 ~]$ curl https://raw.githubusercontent.com/projectcalico/calico/v3.27.0/manifests/custom-resources.yaml -O
  % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload   Total   Spent    Left   Speed
100 810 100 810 0 0 8091 0 --:--:-- --:--:-- --:--:-- 8181
[ec2-user@ip-172-31-93-114 ~]$ kubectl create -f custom-resources.yaml
installation.operator.tigera.io/default created
apiserver.operator.tigera.io/default created
[ec2-user@ip-172-31-93-114 ~]$
```

```
kubectl get nodes ,,,,,,,NOW READY
```

```
[ec2-user@ip-172-31-93-114 ~]$ kubectl get nodes
NAME                                STATUS    ROLES    AGE   VERSION
ip-172-31-81-175.ec2.internal       Ready    <none>    16m   v1.28.7
ip-172-31-91-218.ec2.internal       Ready    <none>    17m   v1.28.7
ip-172-31-93-114.ec2.internal       Ready    control-plane 23m   v1.28.7
[ec2-user@ip-172-31-93-114 ~]$
```

2. Create a NGINX deployment of 3 replicas

vi nginx-deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:latest
        ports:
        - containerPort: 80
```

kubectl apply -f nginx-deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:latest
        ports:
        - containerPort: 80
```

20,0-1 All

kubectl get deployment

```
[ec2-user@ip-172-31-93-114 ~]$ vi nginx-deployment.yaml
[ec2-user@ip-172-31-93-114 ~]$ kubectl apply -f nginx-deployment.yaml
deployment.apps/nginx-deployment created
[ec2-user@ip-172-31-93-114 ~]$ kubectl get deployment
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment    3/3     3             3           53s
[ec2-user@ip-172-31-93-114 ~]$
```